

HW5 (3)

May 1, 2019

```
In [1]: import numpy as np
import pandas as pd
import keras

# plots
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

# processing & model selection
from keras.preprocessing import image
from keras.utils import to_categorical
from sklearn.utils import resample
from sklearn.model_selection import GridSearchCV, train_test_split, StratifiedShuffleSplit

# models
from keras import regularizers
from keras.models import Sequential, Model
from keras.layers import Dense, Activation, BatchNormalization, Dropout, Input, add, Concatenate
from keras.wrappers.scikit_learn import KerasClassifier

# other
import warnings
warnings.filterwarnings('ignore')
np.random.seed(42)

# task3 data loading
import fnmatch
from glob2 import glob
import cv2
```

Using TensorFlow backend.

1 Homework 5

1.1 Task 1

The first task is to run a multilayer perceptron with two hidden layers and relu activations, using the Keras Sequential Interface. We also tune for regularization strength and number of hidden units.

Importing the data:

```
In [2]: df = sns.load_dataset('iris')
        df.head()
```

```
Out[2]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa

Map the target variable:

```
In [3]: print(df.species.unique())
```

```
['setosa' 'versicolor' 'virginica']
```

```
In [0]: mapping = {'setosa': 0, 'versicolor': 1, 'virginica': 2}
        df.species.replace(mapping, inplace = True)
```

Split our dataset into training and test sets.

```
In [0]: X_train, X_test, y_train, y_test = train_test_split(df.drop('species', axis=1), df.species,
```

Encode the target:

```
In [0]: num_species = len(df['species'].unique())

        y_train = to_categorical(y_train, num_species)
        y_test = to_categorical(y_test, num_species)
```

Make our model:

```
In [7]: X_train.shape
```

```
Out[7]: (120, 4)
```

```
In [0]: def make_model(hidden_size, reg_strength, optimizer='adam'):
        """
        Create the model to pass to the Keras Classifier
        """
```

```

model = Sequential([
    Dense(hidden_size, input_shape=(4,), activation='relu',      # first layer
          kernel_regularizer=regularizers.l2(reg_strength)),
    Dense(hidden_size, activation='relu',                        # second layer
          kernel_regularizer=regularizers.l2(reg_strength)),
    Dense(3, activation='softmax')                               # output layer
])

# compile the above model
model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])

return model

```

Before proceeding any further, we should initialize the cross-validation method to be used in the grid search later on.

```
In [0]: sss = StratifiedShuffleSplit(n_splits=3, test_size=0.2, random_state=42)
```

```
In [10]: clf = KerasClassifier(make_model, epochs=10, verbose=0)
```

```

param_grid = {'hidden_size': [32, 64, 128],
              'reg_strength': np.logspace(-3, -1, 5)}

```

```

grid = GridSearchCV(clf, param_grid=param_grid, cv=sss, n_jobs=-1)
grid.fit(X_train, y_train)

```

```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/framework/op_def_registry.py:100:
Instructions for updating:
Colocations handled automatically by placer.
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/math_ops.py:100:
Instructions for updating:
Use tf.cast instead.

```

```

Out[10]: GridSearchCV(cv=StratifiedShuffleSplit(n_splits=3, random_state=42, test_size=0.2,
train_size=None),
error_score='raise-deprecating',
estimator=<keras.wrappers.scikit_learn.KerasClassifier object at 0x7fc5af1445c0>,
fit_params=None, iid='warn', n_jobs=-1,
param_grid={'hidden_size': [32, 64, 128], 'reg_strength': array([0.001, 0.003, 0.01, 0.03, 0.1])},
pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
scoring=None, verbose=0)

```

```

In [11]: print('Best score: {}'.format(np.round(grid.best_score_, 2)))
print('Best params: {}'.format(grid.best_params_))

```

```
Best score: 0.92
```

```
Best params: {'hidden_size': 128, 'reg_strength': 0.001}
```

```
In [12]: score = grid.score(X_test, y_test)
         print('Test score: {}'.format(np.round(score, 2)))
```

Test score: 0.97

1.2 Task 2

Train a multilayer perceptron (fully connected) on the Fashion MNIST dataset. 1. Use 10000 samples from the training set for model selection and to compute learning curves (accuracy vs epochs). 2. Compare the following models and plot their learning curves: - Vanilla Model - Model using dropout - Model using batch normalization and residual connections (but no dropout)

```
In [13]: # loading the data
         (X_train, y_train), (X_test, y_test) = keras.datasets.fashion_mnist.load_data()

Downloading data from http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-labels-
32768/29515 [=====] - 0s 9us/step
Downloading data from http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-images-
26427392/26421880 [=====] - 4s 0us/step
Downloading data from http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/t10k-labels-i
8192/5148 [=====] - 0s 0us/step
Downloading data from http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/t10k-images-i
4423680/4422102 [=====] - 2s 1us/step
```

```
In [14]: print('Training size: {}, Test size: {}'.format(X_train.shape[0], X_test.shape[0]))
         print('Each image has a shape of: {}x{}'.format(X_train.shape[1], X_train.shape[2]))
```

Training size: 60000, Test size: 10000

Each image has a shape of: 28x28

Before proceeding with the modeling, we have to process our data accordingly.

```
In [0]: # reshape X and cast to float
         X_train = X_train.reshape(60000, 784); X_test = X_test.reshape(10000, 784)
         X_train = X_train.astype('float32'); X_test = X_test.astype('float32')

         # scale X
         X_train /= 255; X_test /= 255

         # encode the target accordingly
         n_classes = len(np.unique(y_train))
         y_train = to_categorical(y_train, n_classes); y_test = to_categorical(y_test, n_classes)
```

We also have to get a stratified sample of size 10000 from the training set for model selection.

```
In [0]: X, X_valid, y, y_valid = train_test_split(X_train, y_train, stratify=y_train, random_s

In [0]: # sss = StratifiedShuffleSplit(n_splits=1, test_size=1/6, random_state=6)
         # for a, b in sss.split(X_train, y_train):
         #     X_sample, y_sample = X_train[b], y_train[b]
```

1.2.1 Parameter Tuning

```
In [0]: sss = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=410)
```

Vanilla Model

We start off by tuning our vanilla model. Vanilla models are typically used to describe simple networks with 1 layer, so we'll only look at optimizing the number of hidden units in this case.

```
In [0]: def make_model(hidden_size):
        """
        Create the model to pass to the Keras Classifier
        """

        model = Sequential([
            Dense(hidden_size, input_shape=(784,), activation='relu'), # first layer
            Dense(10, activation='softmax')                             # output layer
        ])

        # compile the above model
        model.compile(optimizer='adam', loss='categorical_crossentropy',
                      metrics=['accuracy'])

        return model
```

```
In [0]: clf = KerasClassifier(make_model, epochs=10, verbose=0)

        param_grid = {'hidden_size': [256, 512]}

        grid = GridSearchCV(clf, param_grid=param_grid, cv=sss, n_jobs=-1).fit(X, y)
```

```
In [21]: print('VANILLA MODEL TUNING:')
          print('Best Score: {}'.format(np.round(grid.best_score_, 2)))
          print('Best Hidden Size: {}'.format(grid.best_params_['hidden_size']))
```

```
VANILLA MODEL TUNING:
Best Score: 0.88
Best Hidden Size: 512
```

```
In [0]: vanilla = grid.best_estimator_
        history_vanilla = vanilla.fit(X, y, epochs=10, validation_data=(X_valid, y_valid))
```

Bigger Model w/ Dropout

We can now proceed with tuning our bigger model with dropout.

```
In [0]: def make_model(hidden_size, hidden_layers, dp_p):
        """
        Create the model to pass to the Keras Classifier
        """
```

```

model = Sequential()

# input layer
model.add(Dense(hidden_size, input_shape=(784,), activation='relu'))
model.add(Dropout(dp_p))

# add as many layers as needed
for i in range(hidden_layers):
    model.add(Dense(hidden_size, activation='relu'))
    model.add(Dropout(dp_p))

# output layer
model.add(Dense(10, activation='softmax'))

# compile the above model
model.compile(optimizer='adam', loss='categorical_crossentropy',
              metrics=['accuracy'])

return model

```

```
In [24]: clf = KerasClassifier(make_model, epochs=10, verbose=0)
```

```

param_grid = {'hidden_size': [256, 512],
              'hidden_layers': [1,2,3],
              'dp_p': np.linspace(0.25, 0.75, 3)}

```

```
grid = GridSearchCV(clf, param_grid=param_grid, cv=sss, n_jobs=-1).fit(X, y)
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:306: *tf.nn.conv2d* is deprecated and will be removed in a future version. Instructions for updating:
Please use ``rate`` instead of ``keep_prob``. Rate should be set to ``rate = 1 - keep_prob``.

```
In [25]: print('DROPOUT MODEL TUNING:')
print('Best Score: {}'.format(np.round(grid.best_score_, 2)))
print('Best Hidden Size: {}'.format(grid.best_params_['hidden_size']))
print('Best Number of Hidden Layers: {}'.format(grid.best_params_['hidden_layers']))
print('Best Dropout Prob: {}'.format(grid.best_params_['dp_p']))
```

DROPOUT MODEL TUNING:

Best Score: 0.88

Best Hidden Size: 512

Best Number of Hidden Layers: 1

Best Dropout Prob: 0.25

```
In [0]: dropout = grid.best_estimator_
history_dropout = dropout.fit(X, y, epochs=10, validation_data=(X_valid, y_valid))
```

Residual Batch Model

Moving on to our final model for this task - we introduce batch normalization and residual connections.

```
In [0]: inputs = Input(shape=(784,))           # input

        x1 = Dense(256, activation='relu')(inputs)  # first layer: 256 units

        x2 = Dense(256, activation='relu')(x1)      # second layer: 256 units
        x2 = BatchNormalization()(x2)              # batch normalization
        x2 = add([x1, x2])                         # residual connection

        x3 = Dense(256, activation='relu')(x2)      # third layer: 256 units
        x3 = BatchNormalization()(x3)              # batch normalization
        x3 = add([x2, x3])                         # residual connection

        x4 = Dense(128, activation='relu')(x3)      # fourth layer: 128 units
        x4 = BatchNormalization()(x4)              # batch normalization

        preds = Dense(10, activation='softmax')(x4)

        resbatch = Model(inputs=inputs, outputs=preds)

In [0]: resbatch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
        history_resbatch = resbatch.fit(X, y, epochs=10, validation_data=(X_valid, y_valid), v


In [29]: for k, h in {'vanilla': history_vanilla,
                    'dropout': history_dropout,
                    'resbatch': history_resbatch}.items():
        print(k.upper() + ' RESULTS:')
        print('Training Accuracy: {}'.format(np.round(h.history['acc'][-1], 2)))
        print('Validation Accuracy: {}'.format(np.round(h.history['val_acc'][-1], 2)))

VANILLA RESULTS:
Training Accuracy: 0.91
Validation Accuracy: 0.9
DROPOUT RESULTS:
Training Accuracy: 0.89
Validation Accuracy: 0.89
RESBATCH RESULTS:
Training Accuracy: 0.9
Validation Accuracy: 0.89
```

1.2.2 Learning Curves

```
In [30]: fig, ax = plt.subplots(1,3, figsize=(18, 4))
        pd.DataFrame(history_vanilla.history).plot(ax=ax[0])
        ax[0].set_title("Vanilla model")
```

```
pd.DataFrame(history_dropout.history).plot(ax=ax[1])
ax[1].set_title("Dropout model")
pd.DataFrame(history_resbatch.history).plot(ax=ax[2])
ax[2].set_title("Batch normalization / Residual connections model")
plt.show()
```



1.3 Task 3

Before proceeding, we should get the data. We have downloaded the dataset onto a Google Drive, where we can unzip the file and load the images.

```
In [3]: from google.colab import drive
drive.mount('/content/drive', force_remount=True)
```

Mounted at /content/drive

```
In [0]: # unzip the new zip file and store all its files into our images directory
!unzip -q -o drive/'My Drive'/'AML HW5'/IDC_regular_ps50_idx5.zip -d images
```

```
In [0]: # load the images
imagePatches = glob('images/**/*.png', recursive=True)

# separate our two classes
class0 = fnmatch.filter(imagePatches, '*class0.png')
class1 = fnmatch.filter(imagePatches, '*class1.png')

# create our target variable
y = []
for img in imagePatches:
    if img in class0:
        y.append(0)
    elif img in class1:
        y.append(1)
```

We turn the images and the target into a dataframe so we can see the class proportions and create a test set.


```
In [0]: images_df = pd.DataFrame()
        images_df['images'] = imagePatches
        images_df['y'] = y
```

```
In [7]: images_df.groupby('y').size()
```

```
Out[7]: y
0      198738
1       78786
dtype: int64
```

Given the imbalanced proportions of our dataset, we implement downsampling.

```
In [0]: images_0 = images_df[images_df.y == 0]  # all our images w/ class 0
        images_1 = images_df[images_df.y == 1]  # all our images w/ class 1

        # downsample the majority class (0)
        images_0 = resample(images_0, replace=False, n_samples=images_1.shape[0], random_state=42)
        # images_1 = resample(images_1, replace=False, n_samples=100000, random_state=42)

        df = pd.concat([images_0, images_1])    # join the two back together
```

```
In [9]: df.head(3)
```

```
Out[9]:
```

	images	y
21948	images/10268/0/10268_idx5_x3501_y1101_class0.png	0
83757	images/14211/0/14211_idx5_x701_y2101_class0.png	0
181091	images/13616/0/13616_idx5_x1551_y701_class0.png	0

```
In [0]: df = df.sample(frac=1).reset_index(drop=True)
```

We can now process our images:

```
In [0]: y = df.y
        X = np.asarray([cv2.resize(cv2.imread(img), (50,50)) for img in df.images])
```

We create a test set and encode/process our data accordingly:

```
In [0]: # encode the target accordingly
        n_classes = len(np.unique(y))                # number of classes in our target
        y = to_categorical(y, n_classes)             # encode training

        # cast X to float and scale
        X = X.astype('float32')
        X /= 255
```

Now that our data is ready, we can proceed with our tasks.

3.1: A model without residual connections.

3.2: Augment the data. How much has this improved the original model?

3.3: Build a deeper model using residual connections. Show that you can build a deep model that would not be able to learn without the residual connections.

3.1 Simple Model

```
In [13]: input_shape = (50, 50, 3)
         simple = Sequential()

         simple.add(Conv2D(32, (2, 2), activation='relu',
                           input_shape=input_shape))
         simple.add(BatchNormalization())
         simple.add(MaxPooling2D(pool_size=(2, 2)))

         simple.add(Conv2D(32, (2, 2), activation='relu'))
         simple.add(BatchNormalization())
         simple.add(MaxPooling2D(pool_size=(2, 2)))

         simple.add(Flatten())
         simple.add(Dense(128, activation='relu'))

         simple.add(Dense(2, activation='softmax'))

         simple.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

Layer 1 ###
Convolutional L

Batch Norm.
Pooling Layer

Layer 2 ###
Conv. Layer
Batch Norm.
Pooling Layer

Layer 3 ###
Flatten the data
64 units dense

Output

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/framework/op_def_registry.py:111: Instructions for updating:
Colocations handled automatically by placer.

```
In [14]: history_simple = simple.fit(X, y, batch_size=128, epochs=10, verbose=1,
                                     validation_split=0.1)
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/math_ops.py:111: Instructions for updating:
Use tf.cast instead.
Train on 141814 samples, validate on 15758 samples
Epoch 1/10
141814/141814 [=====] - 21s 145us/step - loss: 0.4027 - acc: 0.8292 -
Epoch 2/10
141814/141814 [=====] - 17s 121us/step - loss: 0.3528 - acc: 0.8502 -
Epoch 3/10
141814/141814 [=====] - 17s 121us/step - loss: 0.3371 - acc: 0.8576 -
Epoch 4/10
141814/141814 [=====] - 18s 124us/step - loss: 0.3245 - acc: 0.8627 -
Epoch 5/10
141814/141814 [=====] - 17s 121us/step - loss: 0.3112 - acc: 0.8694 -
Epoch 6/10
141814/141814 [=====] - 17s 122us/step - loss: 0.2958 - acc: 0.8767 -
Epoch 7/10

```


141814/141814 [=====] - 17s 122us/step - loss: 0.2799 - acc: 0.8847 -
Epoch 8/10
141814/141814 [=====] - 18s 125us/step - loss: 0.2619 - acc: 0.8914 -
Epoch 9/10
141814/141814 [=====] - 18s 124us/step - loss: 0.2435 - acc: 0.9003 -
Epoch 10/10
141814/141814 [=====] - 17s 122us/step - loss: 0.2255 - acc: 0.9081 -

```

```

In [29]: plt.plot(history_simple.history['acc'], label='train')
plt.plot(history_simple.history['val_acc'], label='test')
plt.title("Simple CNN")
plt.legend()
plt.ylim(0,1)
plt.show()

```



```

In [16]: print('Training Accuracy: {}'.format(np.round(history_simple.history['acc'][-1], 2)))
print('Validation Accuracy: {}'.format(np.round(history_simple.history['val_acc'][-1], 2)))

```

```

Training Accuracy: 0.91
Validation Accuracy: 0.74

```

As we can see, our model has a good training accuracy, but appears to overfit a little. The validation accuracy in particular fluctuates a lot.

3.2 Augmenting the data We're now going to try data augmentation to improve the results of the same CNN - this technique will likely lead our network to overfit less. Here are the options we have:

- Mirroring
- Rotations
- Translations
- Scaling (zooming in/out)

When taking a look at the data, it makes sense to apply all these transformations to the images we have, since they don't have a "normalized" size, orientation, etc. If we were working on the digits dataset, it wouldn't make sense to apply rotations for instance, but here we don't have that problem and we can apply pretty much what we want, since a quick look at the images tells us that they don't seem to be standardized in their scale, orientation, centering, etc.

We're going to do what's called "online" data augmentation, i.e. data augmentation "on the fly", whenever we train a batch. Otherwise, the dataset we'd create would be way too large.

```
In [17]: data_aug = image.ImageDataGenerator(
    horizontal_flip=True,          ## Mirroring (horizontal)
    vertical_flip=True,           ## Mirroring (vertical)
    rotation_range=90,            ## Rotation
    width_shift_range=0.2,        ## Translation (horizontal)
    height_shift_range=0.2,       ## Translation (vertical)
    zoom_range=[1,2],             ## Zooming in
    validation_split=0.1
)

train_generator = data_aug.flow(X, y, batch_size=128, subset='training')
val_generator = data_aug.flow(X, y, batch_size=128, subset='validation')

history_aug = simple.fit_generator(train_generator,
                                   steps_per_epoch=len(train_generator)*4,
                                   # We choose to get 4 times as much data on
                                   # each epoch
                                   epochs=10, verbose=1,
                                   validation_data=val_generator,
                                   validation_steps=len(val_generator))

Epoch 1/10
4432/4432 [=====] - 424s 96ms/step - loss: 0.3368 - acc: 0.8577 - val.
Epoch 2/10
4432/4432 [=====] - 422s 95ms/step - loss: 0.3216 - acc: 0.8653 - val.
Epoch 3/10
4432/4432 [=====] - 423s 96ms/step - loss: 0.3152 - acc: 0.8679 - val.
Epoch 4/10
4432/4432 [=====] - 422s 95ms/step - loss: 0.3100 - acc: 0.8707 - val.
Epoch 5/10
4432/4432 [=====] - 420s 95ms/step - loss: 0.3064 - acc: 0.8721 - val.
Epoch 6/10
```

```


4432/4432 [=====] - 416s 94ms/step - loss: 0.3028 - acc: 0.8738 - val.
Epoch 7/10
4432/4432 [=====] - 416s 94ms/step - loss: 0.3006 - acc: 0.8745 - val.
Epoch 8/10
4432/4432 [=====] - 416s 94ms/step - loss: 0.2972 - acc: 0.8762 - val.
Epoch 9/10
4432/4432 [=====] - 415s 94ms/step - loss: 0.2957 - acc: 0.8765 - val.
Epoch 10/10
4432/4432 [=====] - 427s 96ms/step - loss: 0.2936 - acc: 0.8774 - val.

```

```

In [28]: plt.plot(history_aug.history['acc'], label='train')
plt.plot(history_aug.history['val_acc'], label='test')
plt.title("Simple CNN with augmented data")
plt.legend()
plt.ylim(0,1)
plt.show()

```



```

In [19]: print('Training Accuracy: {}'.format(np.round(history_aug.history['acc'][-1], 2)))
print('Validation Accuracy: {}'.format(np.round(history_aug.history['val_acc'][-1], 2)))

```

```

Training Accuracy: 0.88
Validation Accuracy: 0.86

```

At the expense of a much increased training time, we can see that the same model with augmented data does a little bit better in terms of validation accuracy. The plot also show us that the validation accuracy fluctuates less and generally stays higher. The final accuracy numbers tell us that we indeed overfit less, as the training accuracy has decreased and the validation accuracy is much increased.

3.3 Deeper model Here is our deeper model with residual connections. The architecture is inspired by resnet.

```
In [0]: inputs = Input(shape = (50, 50, 3))

x1 = Conv2D(32, (3, 3), activation='relu')(inputs)
x2 = BatchNormalization()(x1)
x4 = MaxPooling2D(pool_size=(2, 2))(x2)

x5 = Conv2D(32, (3, 3), padding='same',
            activation='relu')(x4)
x6 = BatchNormalization()(x5)

x8 = Conv2D(32, (3, 3), padding='same',
            activation='relu')(x6)
x9 = BatchNormalization()(x8)

x11 = add([x4, x9])

x12 = Conv2D(32, (3, 3), padding='same',
            activation='relu')(x11)
x13 = BatchNormalization()(x12)

x15 = Conv2D(32, (3, 3), padding='same',
            activation='relu')(x13)
x16 = BatchNormalization()(x15)

x80 = add([x11, x16])

x99 = Conv2D(32, (3, 3), padding='same',
            activation='relu')(x80)
x98 = BatchNormalization()(x99)

x97 = Conv2D(32, (3, 3), padding='same',
            activation='relu')(x98)
x96 = BatchNormalization()(x97)

x18 = add([x80, x96])

x19 = Conv2D(64, (3, 3), padding='same',
```

First layer
Conv Layer
Batch Norm.
Pooling Layer

Conv Layer
Batch Norm.

Conv Layer
Batch Norm.

Residual connection

Conv Layer
Batch Norm.

Conv Layer
Batch Norm.

Residual connection

Conv Layer
Batch Norm.

Conv Layer
Batch Norm.

Residual connection

```

        activation='relu')(x18)                                # Conv Layer
x20 = BatchNormalization()(x19)                                # Batch Norm.

x22 = Conv2D(64, (3, 3), padding='same',                      # Conv Layer
            activation='relu')(x20)                            # Batch Norm.
x23 = BatchNormalization()(x22)

x26 = Conv2D(64, (3, 3), padding='same',                      # Conv Layer
            activation='relu')(x23)                            # Batch Norm.
x27 = BatchNormalization()(x26)

x28 = add([x27, x20])                                         # Residual connection

x29 = Conv2D(64, (3, 3), padding='same',                      # Conv Layer
            activation='relu')(x28)                            # Batch Norm.
x30 = BatchNormalization()(x29)

x31 = Conv2D(64, (3, 3), padding='same',                      # Conv Layer
            activation='relu')(x30)                            # Batch Norm.
x32 = BatchNormalization()(x31)

x40 = add([x32, x28])                                         # Residual connection

x41 = Conv2D(64, (3, 3), padding='same',                      # Conv Layer
            activation='relu')(x40)                            # Batch Norm.
x42 = BatchNormalization()(x41)

x43 = Conv2D(64, (3, 3), padding='same',                      # Conv Layer
            activation='relu')(x42)                            # Batch Norm.
x44 = BatchNormalization()(x43)

x45 = add([x44, x40])                                         # Residual connection

x46 = Conv2D(64, (3, 3), padding='same',                      # Conv Layer
            activation='relu')(x45)                            # Batch Norm.
x47 = BatchNormalization()(x46)

x48 = Conv2D(64, (3, 3), padding='same',                      # Conv Layer
            activation='relu')(x47)                            # Batch Norm.
x49 = BatchNormalization()(x48)

x33 = add([x45, x49])                                         # Residual connection

x34 = Flatten()(x33)                                         ### Last layer ###
x35 = Dense(64, activation='relu')(x34)                       # Flatten the data
                                                                # 64 units dense

preds = Dense(2, activation='softmax')(x35)                  ### Output ###

```

```

deep_res = Model(inputs=inputs, outputs=preds)
deep_res.compile(optimizer='adam', loss='categorical_crossentropy',
                 metrics=['accuracy'])

```

And here is the same model without residual connections.

```

In [0]: input_shape = (50, 50, 3)
        deep_nores = Sequential()

        deep_nores.add(Conv2D(32, (3, 3), activation='relu',
                               input_shape=input_shape, padding='same'))
        deep_nores.add(BatchNormalization())
        deep_nores.add(MaxPooling2D(pool_size=(2, 2)))

        for i in range(6):
            deep_nores.add(Conv2D(32, (3, 3), activation='relu',
                                   padding='same'))
            deep_nores.add(BatchNormalization())

        for i in range(9):
            deep_nores.add(Conv2D(32, (3, 3), activation='relu',
                                   padding='same'))
            deep_nores.add(BatchNormalization())

        deep_nores.add(Flatten())
        deep_nores.add(Dense(64, activation='relu'))

        deep_nores.add(Dense(2, activation='softmax'))

        deep_nores.compile(optimizer='adam', loss='categorical_crossentropy',
                           metrics=['accuracy'])

```

Layer 1
Convolutional L

Batch Norm.
Pooling Layer

Layers 2-7

Conv. Layer
Batch Norm.

Layers 8-16

Conv. Layer
Batch Norm.

Layer 17
Flatten the data
64 units dense

Output

Let's now train them.

```

In [45]: history_res = deep_res.fit(X, y, batch_size=128, epochs=10, verbose=1,
                                     validation_split=0.1)

```

Train on 141814 samples, validate on 15758 samples

Epoch 1/10

141814/141814 [=====] - 110s 779us/step - loss: 0.4136 - acc: 0.8314

Epoch 2/10

141814/141814 [=====] - 103s 726us/step - loss: 0.3477 - acc: 0.8540

Epoch 3/10

141814/141814 [=====] - 103s 726us/step - loss: 0.3315 - acc: 0.8606

Epoch 4/10

141814/141814 [=====] - 103s 724us/step - loss: 0.3180 - acc: 0.8671


```

Epoch 5/10
141814/141814 [=====] - 103s 725us/step - loss: 0.3100 - acc: 0.8710
Epoch 6/10
141814/141814 [=====] - 103s 724us/step - loss: 0.2989 - acc: 0.8752
Epoch 7/10
141814/141814 [=====] - 103s 725us/step - loss: 0.2867 - acc: 0.8798
Epoch 8/10
141814/141814 [=====] - 103s 726us/step - loss: 0.2739 - acc: 0.8846
Epoch 9/10
141814/141814 [=====] - 103s 724us/step - loss: 0.2615 - acc: 0.8888
Epoch 10/10
141814/141814 [=====] - 103s 726us/step - loss: 0.2410 - acc: 0.8970

```

```

In [46]: history_nores = deep_nores.fit(X, y, batch_size=128, epochs=10, verbose=1,
                                         validation_split=0.1)

```

Train on 141814 samples, validate on 15758 samples

```

Epoch 1/10
141814/141814 [=====] - 78s 549us/step - loss: 0.4050 - acc: 0.8263
Epoch 2/10
141814/141814 [=====] - 71s 497us/step - loss: 0.3597 - acc: 0.8471
Epoch 3/10
141814/141814 [=====] - 70s 496us/step - loss: 0.3447 - acc: 0.8542
Epoch 4/10
141814/141814 [=====] - 71s 499us/step - loss: 0.3364 - acc: 0.8586
Epoch 5/10
141814/141814 [=====] - 71s 499us/step - loss: 0.3318 - acc: 0.8614
Epoch 6/10
141814/141814 [=====] - 70s 497us/step - loss: 0.3253 - acc: 0.8638
Epoch 7/10
141814/141814 [=====] - 71s 498us/step - loss: 0.3199 - acc: 0.8669
Epoch 8/10
141814/141814 [=====] - 71s 498us/step - loss: 0.3140 - acc: 0.8693
Epoch 9/10
141814/141814 [=====] - 71s 498us/step - loss: 0.3082 - acc: 0.8726
Epoch 10/10
141814/141814 [=====] - 71s 498us/step - loss: 0.3025 - acc: 0.8742

```

```

In [47]: fig, ax = plt.subplots(1,2, figsize=(12, 4))
         ax[0].plot(history_res.history['acc'], label='train')
         ax[0].plot(history_res.history['val_acc'], label='test')
         ax[0].set_title("Deep CNN with residual connections")
         ax[0].set_ylim(0,1)
         ax[0].legend()

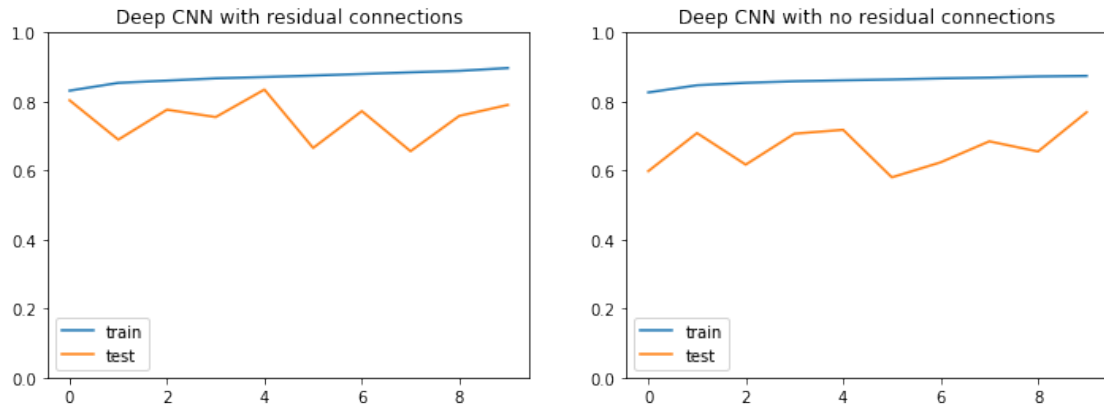
         ax[1].plot(history_nores.history['acc'], label='train')

```

```

ax[1].plot(history_nores.history['val_acc'], label='test')
ax[1].set_title("Deep CNN with no residual connections")
ax[1].set_ylim(0,1)
ax[1].legend()
plt.show()

```



As we can see, removing the residual connections makes the performance decrease - although not by much. We expect the difference in performance to be even higher with an even deeper model.